

Paper title

KV Cache is 1 Bit Per Channel: Efficient Large Language Model Inference with Coupled Quantiza

Exploring KV Cache Quantization in Multimodal Large Language Model Inference

AirCache: Activating Inter-Modal Relevancy KV Cache Compression for Efficient Large Vision-Lar

Towards Efficient Large Language Model Serving: A Survey on System-Aware KV Cache Optimiz

Towards Efficient Large Language Model Serving: A Survey on System-Aware KV Cache Optimiz

Accelerating Chatbot Inference with vLLM: Evaluating the Efficiency of PagedAttention

VLLM Guided Human-Like Guidance Navigation Generation

Enhancing Flood Prediction Reliability with RAFT-VLLM: A Retrieval Augmented and Fine-Tuned

Scaling AI Infrastructure: From Recommendation Engines to LLM Deployment with PagedAttenti

FlexEdit: Marrying Free-Shape Masks to VLLM for Flexible Image Editing

H2O: Heavy-Hitter Oracle for Efficient Generative Inference of Large Language Models

Evolution of Cache Replacement Policies to Track Heavy-Hitter Flows

Hierarchical Heavy Hitter Mining on Streams

Heavy Hitter

Evolution of cache replacement policies to track heavy-hitter flows

Topic	DOI
KV cache	10.52202/079017-0109
KV cache	10.1109/lca.2025.3646170
KV cache	10.1109/iccv51701.2025.02221
KV cache	10.36227/techrxiv.176046306.66521015/v2
KV cache	10.36227/techrxiv.176046306.66521015/v3
PagedAttention	10.1109/qpain66474.2025.11171881
PagedAttention	10.5220/0013191100003912
PagedAttention	10.1016/j.procs.2025.09.439
PagedAttention	10.32996/jcsts.2025.7.4.39
PagedAttention	10.2139/ssrn.5933540
KV eviction	10.52202/075280-1506
KV eviction	10.1007/978-3-642-19260-9_3
KV eviction	10.1007/springerreference_63323
KV eviction	10.5040/9780571343911.ch-003
KV eviction	10.1145/1872007.1872046